

A feature-based taxonomy of coordination approaches

Tim Kräuter¹, Julien Deantoni², Adrian Rutle¹, Harald König^{3,1}, and
Yngve Lamo¹

¹ Western Norway University of Applied Sciences, Bergen, Norway
`tkra@hvl.no`, `aru@hvl.no`, `yla@hvl.no`

² University Cote d’Azur, Sophia Antipolis, France
`julien.deantoni@univ-cotedazur.fr`

³ University of Applied Sciences, FHDW, Hanover, Germany
`harald.koenig@fhdw.de`

Abstract. Complex domains necessitate the utilization of multiple interacting software systems. Various categories of coordination approaches, such as coordination languages, co-simulation approaches, architecture description languages, and coordination frameworks, have been proposed to ensure seamless integration of these systems. In this paper, we present a comprehensive feature-based taxonomy of these categories, which have previously only been studied in isolation. The taxonomy uncovers common and unique features across the coordination approaches. It can be used to make informed decisions about the choice of coordination approaches for specific use cases.

Keywords: Co-simulation · Coordination language · ADL · Coordination framework · Taxonomy · Feature model

1 Introduction

Complex domains require multiple interacting software systems to meet their demands. Coordination approaches are necessary to ensure the effective integration of these systems. The study of coordination approaches appears to be fragmented, with disparate investigations conducted on different abstraction levels, pursuing varied goals, and utilizing diverse terminology, resulting in different categories of approaches, namely *coordination languages* ([43]), *Co-simulation approaches* ([19]), *architecture description languages* ([14]), and *coordination frameworks* ([26,53]). To the best of our knowledge, evaluations of coordination approaches have been limited to their respective communities. This isolation hinders a comprehensive understanding of coordinating software systems, leading to the development of the same features independently with limited reuse.

First, we introduce the different categories of coordination approaches to paint a picture of the different communities (section 2) and discuss our methodology in section 3. Our main contribution is a comprehensive taxonomy, represented as a *feature model* [24] to categorize coordination approaches within

a unified framework ((section 4)). The taxonomy allows the comparison of coordination approaches across the different categories. Moreover, we apply the taxonomy to evaluate 17 distinct approaches spanning the aforementioned categories (section 5) and publish our findings in a public dataset [51]. Finally, we discuss our findings in section 6 before concluding in section 7, meaning we identify the typical features for each category, provide general insights, and cluster the approaches regarding their similarity.

2 Categories of coordination approaches

Figure 1 gives an overview of the different categories of coordination approaches. The categories of coordination approaches broadly operate on different abstraction levels, namely the *execution*, *model*, and *language* levels. We will explain each abstraction level and the accompanying coordination approaches from the bottom up. A detailed explanation of each category is given in the following subsections. In general, the right abstraction level to address coordination depends on the specific use case and goals at hand, meaning approaches on each abstraction level have their advantages and disadvantages.

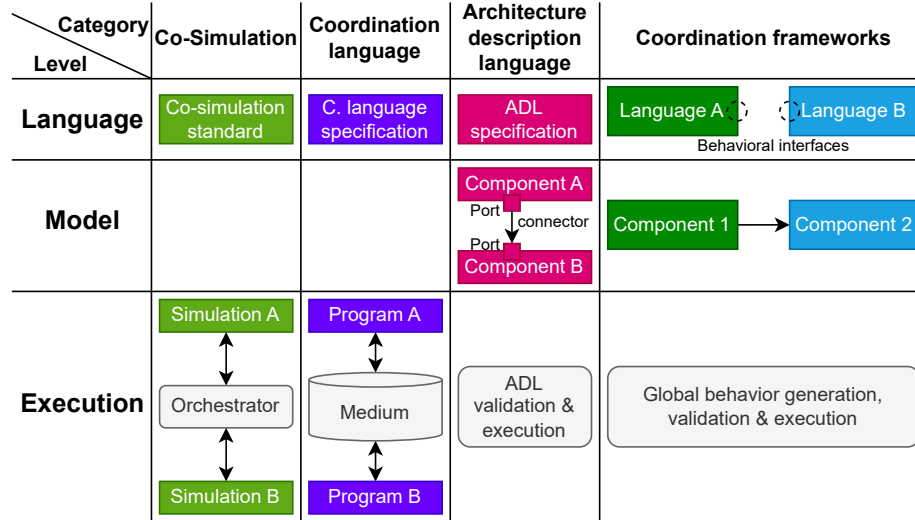


Fig. 1. Overview of coordination approaches

Execution level: Approaches on the *execution level* aim to deal with coordination by interfacing directly with executable binaries or source code. We identified two distinct categories on this level: co-simulation approaches and coordination languages, described in detail in subsection 2.1 and subsection 2.2.

Model level: *Architecture Description Languages (ADLs)* operate on the *model* level, i.e., each *component* is given by a behavioral model, for example, a process algebra. ADLs are described in detail in subsection 2.3.

Language level: *Coordination frameworks* operate on the language level since they allow coordination of models conforming to *different* behavioral languages. They are described in detail in subsection 2.4.

2.1 Co-simulation approaches

Julien will make a proposition.

A *co-simulation approach* composes the simulation of multiple systems into a global simulation using an *orchestrator* [19]. A system simulation also called a simulation unit (SU), is typically given by an executable binary conforming to a predefined interface.

2.2 Coordination languages

Many different styles of coordination languages exist, which achieve coordination between systems differently. In this context, a system is an executable program of varying size, but we also use the term system to describe components in an ADL or behavioral models in coordination frameworks.

Coordination languages such as Linda [12] provide a *set of operations* to enrich a host programming language with coordination capabilities. The operations in Linda are used to read from and write to a virtual shared memory, which serves as a *global* communication medium. Such coordination languages are often called *tuple-based* because the shared memory contains sequences of elements, i.e., tuples. Ttuple-based coordination languages such as Linda are surveyed in [47,38,40].

Other coordination languages such as Manifold [4,44] and REO [6] do not rely on a global communication medium but rather use *channels* to connect systems locally. Each system then uses I/O operations to interact with connected channels without knowing who has sent or will receive data through the channel. These coordination languages share many ideas with Architecture Description Languages (ADLs) discussed in the next section.

Finally, coordination languages might also use communication models such as the actor or reactor [30] model. For example, the Lingua Franca [30,31] coordination language is described using the reactor model, meaning one defines coordination by describing how each system reacts to incoming events. A reaction to an event in one system might trigger reactions in connected systems. To encode reactions, the user can choose an existing programming language, such as C, C++, Python, TypeScript, and Rust. Lingua Franca aims to simplify developing multi-threaded applications using the reactor model to guarantee *determinism* and provide *built-in timing semantics*.

2.3 Architecture description languages

Architecture Description Languages (ADLs) aim to describe the structure of systems, allowing developers to focus on high-level system components and their connections rather than lines of source code [14,36,37]. Many different ADLs have been proposed in the academic literature and by the industry [36,54]. Nevertheless, clearly defining ADLs is challenging due to overlap with general-purpose modeling languages [14]. This distinction is deeply studied in [36].

The three building blocks of ADLs are defined as (1) *components*, (2) *connectors*, (3) *architectural configuration* [36,37]. A *component* is a unit of computation or data repository [36]. Components vary in size, ranging from representing individual services to entire systems. In this paper, we use the more general term system when we are not in the immediate context of ADLs.

Connectors serve as architectural elements to model interactions between components and the regulations that oversee those interactions [36]. A difference to components is that connectors must not be implemented as distinct entities such as message brokers but can also represent shared variables or links between applications realized by client-server protocols [36].

Architectural configuration, also known as topology, represents the structural arrangement of components and connectors in a connected graph, defining the overall architecture [36]. This structure determines if the combined semantics satisfy the desired behavior. Verification relies heavily on specifying components and connectors. For instance, one common application ensures the architectural configuration is free from deadlock and starvation.

The first ADLs were developed in the 90s and use process algebras such as Communicating Sequential Processes (CSP), Calculus of Communicating Systems (CCS), and π -calculus [41]. For example, Wright uses CSP [2] while Darwin uses π -calculus [32]. More recent ADLs, such as MontiArc [21], use automata to define the behavior of components.

However, to the best of our knowledge, no ADL supports heterogeneous components such as the coordination frameworks discussed in the next section [36]. Furthermore, despite the creation of numerous ADLs in the literature, they are not mainstream, i.e., not often used by practitioners in the industry [14,54,42,41,35].

2.4 Coordination frameworks

Julien will make a proposition, which should include somewhere:

We introduce the term *coordination framework* for coordination approaches operating on the language level. Coordination frameworks allow coordination of models conforming to *different* behavioral languages using the concept of *behavioral interfaces*. They allow modelers to utilize different languages for varying aspects of the system. In this paper, we will refer to our approach outlined in [26] as the *Behavioral Coordination Language* (BCoorLang) to enhance clarity in communication.

3 Methodology

First, we conducted a *meta-analysis* (tertiary study) on surveys (secondary studies) within each coordination category. We found the following secondary studies for coordination languages ([43,20,47]), ADLs ([14,36,23,41,33]), and co-simulation approaches ([19,50,22]). We did not find surveys on coordination frameworks since the term is not yet commonly used. We used Google Scholar to locate these papers, utilizing the keywords “survey”, “feature model”, “taxonomy”, and for each category, “coordination language”, “architecture description language”, or “co-simulation” accordingly. We state our exact search queries in [51]. We then collected the first five pages of papers and filtered them by title, abstract, and content to arrive at the previously listed papers.

Second, we investigated individual approaches in each category to construct our taxonomy. Due to the immense number of approaches and limited time, we focus on the 17 approaches listed in section 5. Our selection of these approaches is based on a mix of different criteria: distribution across categories, expert knowledge of the authors regarding relevance, and heterogeneity inside categories. For example, we chose heterogeneous coordination languages, namely Linda, Reo, and Lingua Franca, which employ different coordination models, namely tuple spaces, channels, and reactors.

Finally, we refined our taxonomy by examining the secondary studies and individual coordination approaches until it reached a stable state described in the next section.

4 Taxonomy

We represent our taxonomy as a *feature model* [24]. We aim to comprehensively understand the coordination landscape and compare coordination approaches from different categories. To achieve this goal, we first identify the coordination *foundation*: coordination mechanism, syntax, semantics, and degree of formality. Second, we investigate the coordination approaches *goal* such as simulation, execution, and formal validation. Third, we study the approaches’ *properties*, such as domain, execution, specification, system transparency, and system languages.

The top level of the feature model is shown in Figure 2. A coordination approach has the three previously discussed features: *Foundation*, *Goal*, and *Properties*, which we depict and describe in detail in the following sections. A complete diagram of the feature model is contained in [51].

4.1 Foundation

Figure 3 shows the **Foundation** feature, which consists of **Syntax**, **Semantics**, and **Mechanism**.

Syntax The coordination syntax of the elements on which to apply coordination can either be *explicit* or *implicit*. An implicit syntax is usually based on the

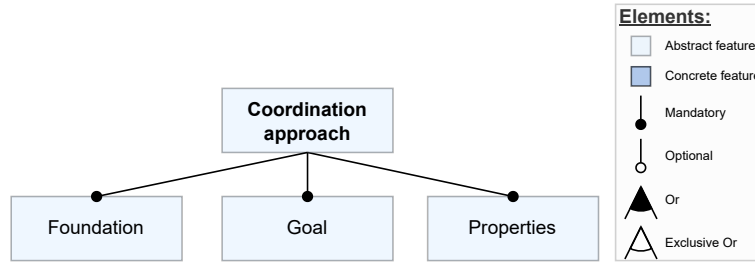


Fig. 2. Feature model overview

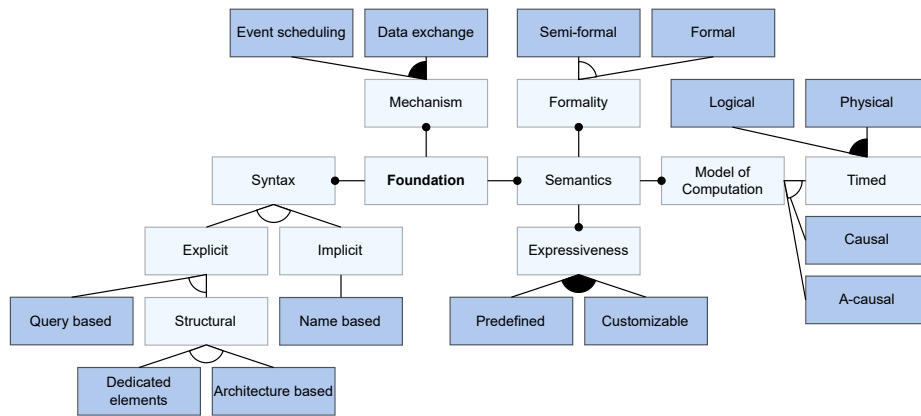


Fig. 3. Foundation feature

names of elements. For example, one could synchronize the firing of transitions from multiple state machines if they have identical names.

Explicit syntax is either *query based* or *structural*. For example, in BCOoL [52,53], one can write queries to define which elements should coordinate.

A structural definition of coordination uses *dedicated elements* or is *architecture based*. For example, interactions in BCoorLang are specified by picking dedicated elements that should be coordinated, while coordination in ADLs depends on the architectural configuration [36], i.e., how ports of systems are connected.

Mechanism The main mechanism driving coordination can be *event scheduling*, *data exchange*, or a mix of both. In the literature, these mechanisms are sometimes called control/event-driven or data-driven, respectively [43,52].

First, *event scheduling* synchronizes events or imposes a specific event order, i.e., order of state changes in different systems. Usually, event scheduling is used when systems are given by behavioral models since there might not be a clear concept of an event in a programming language. For example, in BCOoL [53], one can define relationships between events, such as *happens before* or *synchronize*.

Data-exchange means systems communicate by exchanging data. For example, in the ADL MontiArc [21], systems send data from the output to the input port. Data can also be exposed as variables, which are then coupled with variables in other systems, for example, to be the same. However, data exchange might also happen between systems during the synchronization of events, which is why both coordination mechanisms can be supported simultaneously.

Semantics The semantics feature consists of *expressiveness*, *model of computation*, and *formality*. The expressiveness of coordination can be *predefined*, i.e., a user has a fixed set of coordination mechanisms, for example, a fixed set of connectors with predefined behavior. Expressiveness is *customizable* if one can define the semantics of coordination, typically by using a dedicated language. For example, in BCOoL, one can employ the Clock Constraints Specification Language (CCSL) [3] to define new coordination operators besides the predefined ones [52,53].

Coordination approaches utilize different *models of computation* (MoC). A MoC is a set of rules that governs the timed, possibly concurrent execution and coordination of systems [45]. We categorize MoCs along three types: *timed*, *causal*, *a-causal*. A *timed* MoC can use *physical* or *logical* time. The categorization of *causal* and *a-causal* MoCs is also present in the co-simulation field [19]. The causal MoC means that the coordination approach facilitates event and data exchange between different systems such that one system can cause reactions, i.e., state changes in other systems. For example, a system based on a state machine raises an event, including data, which causes a state change and possible other reactions in a different system.

An a-causal MoC is necessary to deal with related differential equations [27]. For example, a variable in one equation should equal a variable in a different equation. In this case, there is no cause and effect, such as in the causal MoC, but rather, the variables have an a-causal relationship since they should be equal at any time.

Logical time means that coordination approaches define when events/state changes happen. Typically, one defines that events take place simultaneously or that one event causes another event immediately or with a fixed delay. For example, a delay could be 50 milliseconds of logical time, corresponding to time in a simulation that can run arbitrarily fast.

In contrast, *physical* time describes time passing in the physical world. It is also referred to as *wall-clock* time [19]. For example, one can define that an event should happen 50 milliseconds after another. In Lingua Franca [30], logical time “chases” physical time, i.e., logical time tries to match the physical time provided by the execution platform.

Formality of a coordination approach is either *formal* or *semi-formal*. Semi-formal approaches are directly implemented in a general-purpose programming language, often with a guiding formalism in mind but without strict adherence to that formalism.

Formal coordination approaches are implemented in a *formal language*, which can then be run on an execution engine implemented for that formalism. For ex-

ample, the coordination framework BCoorLang [26] uses graph or term-rewriting as formal languages to coordinate heterogeneous systems in a centralized manner. It relies on the graph-transformation tool Groove [46] or the term-rewriting tool Maude for execution.

4.2 Goal

We identified three different *goals* of coordination approaches: **simulation**, **execution**, and **formal validation**, see Figure 4. A coordination approach can have one or multiple goals, which we will now describe in detail.




Fig. 4. Goal feature

Simulation The most common goal of coordination approaches is a coordinated simulation of multiple systems or behavioral models representing an entire system. A coordinated simulation is useful because the composition of multiple systems can lead to unexpected behavior, often called *emergent behavior* [16]. Emergent behavior can be dealt with early in development before implementing each system and its interactions. Co-simulation approaches such as DAC-COSIM [18,15] or MECSYCO (Multi-agent Environment for ComplexSYstem CO-simulation) [10,11] are examples of approaches that have coordinated simulation as their goal. Simulation is also a goal of coordination frameworks such as BCoorLang [26] and BCOol [53]. Even recent ADLs such as MontiArc [21] provide simulation capabilities.

Execution In contrast to coordinated simulation, which often only aims to find errors during development, some coordination approaches aim to provide a coordination infrastructure for system execution upon deployment. For example, Linda provides a virtual shared memory with read and write operations to coordinate systems.

Furthermore, the recently developed polyglot coordination language *Lingua Franca* provides determinism guarantees during distributed execution [31]. Thus, it eliminates typical coordination concerns faced during coordinating concurrent *execution*.

Formal validation Another goal of coordination approaches, especially for safety-critical systems, is formal validation of the coordinated system.

We encountered coordination approaches that validate *predefined properties*, while others support *custom properties*. For example, the ADL Wright [2] can

check *deadlock freedom* of system interactions. The coordination framework for behavioral consistency [26] supports custom properties written in temporal logic such as Linear Temporal Logic (LTL) and Computational Tree Logic (CTL).

4.3 Properties

We identified several reoccurring properties when analyzing coordination approaches. Figure 5 depicts the resulting *properties* feature, which we will now explain in detail.

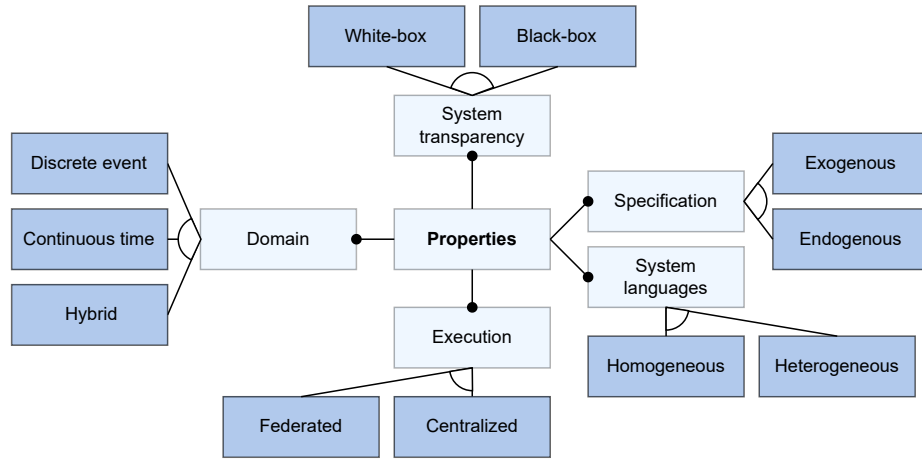


Fig. 5. Properties feature

Domain Coordination approaches operate in different domains like the domains of co-simulation approaches described in [19]. First, if a coordination approach is part of the *Discrete event* (DE) domain, it allows discrete but not continuous state changes [19]. Typically, coordination frameworks and ADLs operate in the DE domain since variables change their values discontinuously during execution, i.e., immediately change between states.

Second, coordination approaches in the *continuous time* (CT) domain support systems to have a state that evolves continuously over time. Often, systems are given by differential equations, which induce continuous variable changes, such as physical systems involving springs and dampeners, see [19].

Third, the *hybrid* domain is a mix of the DE and CT domains that occur, for example, when simulating cyber-physical systems. In a hybrid domain, systems should coordinate where some change their state discretely (DE domain) while others evolve continuously over time (CT domain). Two strategies deal with hybrid systems: either adapt the systems from the CT to the DE domain or vice versa [19].

System transparency Some coordination approaches require constituent systems to be entirely transparent (*white-box*), while others only require a fixed interface (*black-box*). Usually, black-box approaches are needed to protect *intellectual property* (IP) when different companies want to work together without sharing all the details concerning their systems.

For example, DACCOSIM only requires each system to conform to the FMI specification for Co-Simulation.

Specification We distinguish between *endogenous* and *exogenous* specification of coordination, as described for coordination languages in [5]. For example, when using the coordination language Linda [12], one uses coordination-specific operations such as *out* and *in* inside system programs to write to and read from the shared tuple space, respectively. Thus, Linda is *endogenous* since one has to modify the behavior of the systems to add coordination.

On the other hand, *exogenous* coordination approaches keep coordination aspects outside of systems. However, systems must still be created with coordination in mind to expose possible coordination points. For example, BCOol [53] uses language behavioral interfaces to define events, which can then be used to specify *exogenous* coordination rules outside the systems. Endogenous coordination approaches tend to mix computation and coordination [5].

System languages Coordination approaches typically only support *homogeneous* systems. For example, coordination languages, such as Lingua Franca [30], currently only support coordination if each system is given in the same programming language.

However, coordination frameworks allow *heterogeneous* systems, i.e., systems specified in different behavioral modeling languages. For example, BCoorLang [26] supports all modeling languages where one can clearly define *state structure* and which *elements can change state* during execution.

Execution A coordination approach is realized using either a *centralized* or *federated* execution. In a *centralized* execution, all systems are executed by one engine, which enforces coordination. For example, BCorrLang composes the behavior of all system specifications into a global behavior specification, which is then executed by a graph transformation or term-rewriting tool.

On the other hand, in a *federated execution*, systems run independently and only coordinate using a shared infrastructure through predefined connections. Federated execution can still happen on the same physical machine, i.e., it does not mean that systems must be *distributed* across different physical machines.

For example, the coordination language Lingua Franca [30] supports federated execution on one machine by providing a runtime infrastructure. DACCOSIM [18] goes one step further and allows a *distributed* federated execution.

5 Application of the feature model

Table 1 shows the application of our feature model to four different coordination approaches. Each approach comes from a different category: co-simulation (*DAC-*

COSIM [18,15]), coordination language (*Linda* [12,13]), ADL (*MontiArc* [21]), and coordination framework (*BCoorLang* [25,26]).

Table 1. Approach classification

Approach	DACCOSIM	Linda	MontiArc	BCoorLang
Feature				
Foundation				
Syntax	Architecture based	Name based	Architecture based	Dedicated elements
Mechanism	Data-Exchange	Data-Exchange	Data-Exchange	Event-Scheduling
Expressiveness	Predefined	Predefined	Predefined	Predefined
MoC	Causal	Causal	Logical, Causal	Logical
Formality	Semi-formal	Semi-formal	Semi-formal	Formal
Goal				
Simulation	+	-	+	+
Execution	-	+	-	-
Formal verification	-	-	-	Custom properties
Properties				
Domain	Hybrid	Discrete event	Discrete event	Discrete event
Comp. transparency	Black-box	White-box	White-box	White-box
Specification	Exogenous	Endogenous	Endogenous	Exogenous
Comp. languages	Homogeneous	Homogeneous	Homogeneous	Heterogeneous
Execution	Federated	Federated	Centralized	Centralized

In addition, we classified the following approaches: Ptolemy [16,45], Wright [2,1], CommUnity [17,39], Metropolis [7], UMoC++ [34], Lingua Franca [30,31], Reo [6], BIP [9,8], MECSYCO [11,10], CoSim20 [29], BCOoL [53,52], Manifold [4,44], and ForSyDe [48,49]. Due to space constraints, we cannot show all classifications, but the full data set is available in the artifacts of this paper [51].

In Figure 6, we compare BCoorLang and MontiArc with Linda on the left and DACCOSIM on the right. The figure shows Venn diagrams comparing the feature sets of each approach, where each colored part states which features are contained. In addition, the number in the intersection states the amount of common features.

For example, when looking at the Venn diagram on the left in Figure 6, the 1 in the circle part at the bottom describes that MontiArc has one *original* feature (Definition:Architecture based), while the 3 in the intersection of all approaches in the middle declares that they all have three features in *common* (Expressiveness:Predefined, Domain:Discrete event, System transparency:White-box), see Table 1.

The Venn diagram on the left in Figure 6 shows that Linda has three original features compared to BCoorLang and MontiArc: Goal:Execution, Definition:Name based, and Execution:Federated. Furthermore, one can see that coordination frameworks, i.e., BCoorLang (the diagram is similar for BCOoL [53,52])

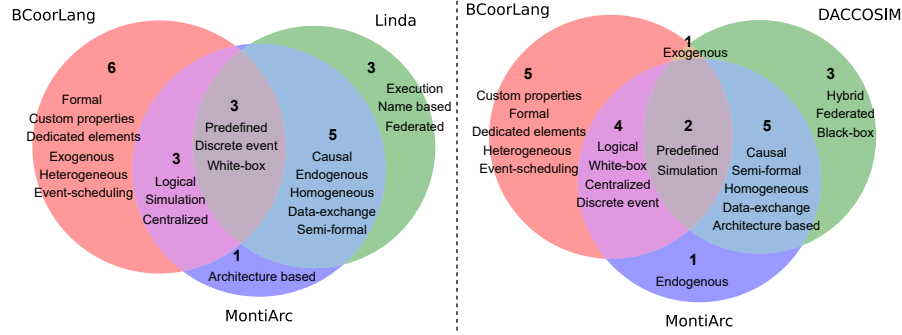


Fig. 6. Feature comparison: BCoorLang, MontiArc, and Linda/DACCOSIM

greatly differ from ADLs and coordination languages. Coordination frameworks have multiple original features, such as `System languages:Heterogeneous` and `Formal verification:Custom properties`.

Similarly, on the right in Figure 6 BCoorLang and MontiArc are compared with the co-simulation approach DACCOSIM. One can see that BCoorLang and DACCOSIM do not have many features in common, probably due to their different goals as a coordination approach. Co-simulation approaches like DACCOSIM have original features, such as `Execution:Federated`, `Domain:Hybrid`, and `System transparency:Black-box`.

Venn diagrams give a good approximation of the similarity between three specific coordination approaches. However, they do not tell the whole story since each feature is given equal weight, which is not always adequate. For example, a difference in the *domain* feature is more profound than most other differences. We stick to this simple weighting throughout this paper since the importance of features can be subjective. However, one can adapt the analysis using the provided data and scripts [51].

6 Findings

In this section, we describe the findings of applying our taxonomy to the different coordination approaches. We start by summarizing the typical features of each coordination category before stating the general insights we have attained. Finally, we cluster the approaches based on their feature similarity and discuss the results.

6.1 Typical features

We will describe the common feature sets we found for each category.

ADLs primarily have *simulation* as their goal, while a small subset allows predefined properties to be verified. Unsurprisingly, the coordination syntax of

all ADLs is architecture-based, using components with ports and connectors between them. All ADLs operate in the discrete event domain. Furthermore, components in ADLs are white-box and homogeneous, and coordination specification is usually exogenous, i.e., part of the connectors. An exception is MontiArc, where ports can be marked as *synchronized* only allowing synchronous interactions through this port, meaning coordination specification is partly present inside components.

Coordination languages can have simulation and execution as their goal while operating in the discrete event domain. Surprisingly, coordination syntax is also most likely architecture-based, i.e., the idea of subsystems, ports, and connectors/channels between them is also used in most coordination languages. In addition, systems in coordination languages are also white-box and homogenous but have both endogenous and exogenous coordination specifications. Notably, some coordination languages offer a federated execution not supported by ADLs.

Unsurprisingly, all **co-simulation** approaches have only simulation as their goal. Similarly to ADLs and coordination languages, coordination syntax is architecture-based. However, compared to the other categories, they operate in the hybrid domain, facilitating data exchange between black-box systems. Systems can be homogeneous or heterogeneous by wrapping them in specific formalisms such as DEVS in the case of MECSYCO. Typically, execution is federated, while some approaches, such as DACCOSIM, even allow for a distributed execution.

Coordination frameworks have simulation and formal validation as their goal. Coordination syntax is diverse but usually uses event scheduling in the discrete event domain. Furthermore, coordination specification is exogenous to support heterogeneous systems, which are white-box. In contrast to coordination languages and co-simulation approaches, execution is centralized. Interestingly, coordination frameworks such as BCoorLang allow formal validation of custom properties, a rare feature.

6.2 General insights

Architecture-based coordination syntax is widely used. Remarkably, coordination approaches from all categories use concepts such as components/subsystems, ports, and connectors to express coordination. The specific coordination details vary between different approaches, but there appears to be broad consensus on the importance of clearly representing the system architecture for effective coordination.

Formal validation is rarely supported. Most of the studied coordination approaches have simulation as their goal, while formal validation is uncommon. This could be attributed to the problem of formally representing heterogeneous systems engaged in coordination and the rapid growth of state spaces when multiple systems are involved.

Coordination frameworks do not support Data-Exchange. All coordination frameworks besides Ptolemy support Event-Scheduling but not Data-Exchange as a coordination mechanism. This is likely because data exchange

is difficult to achieve when heterogeneous systems are allowed. However, data exchange is omnipotent in real-world systems, so it is the main coordination mechanism used in co-simulation approaches and coordination languages. Thus, for coordination frameworks to be employed in real-world scenarios, they must support data exchange.

6.3 Feature clustering analysis

We further analyze the classification of the approaches mentioned in section 5 to find out how they compare across and inside the four categories. We clustered the approaches to see similarities and differences by just looking at the feature data of each approach, not considering to which category they belong. The data and scripts to reproduce our analysis are available in [51]. We apply a standard clustering algorithm using Jaccard distance to measure the similarity of the feature sets. Jaccard distance is the one-complement of Jaccard similarity for two sets, defined as the ratio of their intersection to their union [28]. It is a widely used and reasonable dissimilarity measure for sets [28].

Figure 7 shows a scatter plot with approximate positions of each approach derived solely from the distances between them. It provides a bird’s-eye view of the similarity or dissimilarity between the classified coordination approaches, in contrast to the previous Venn diagrams, which only compare three approaches. Furthermore, it highlights the clustering results by coloring each data point. We selected clustering parameters to reduce the number of unclustered approaches while ensuring that we do not consolidate everything into just one cluster.

The clustering algorithm finds three clusters, including all but one approach. **Cluster 1** contains BCoorLang and BCOol, only coordination frameworks, while **cluster 2** consists of a mix of coordination languages and ADLs, such as Linda and Lingua Franca, as well as Wright and MontiArc, respectively. **Cluster 3** includes three co-simulation approaches, but no cluster contains Ptolemy.

We interpret cluster 1 as representing the coordination framework category. Similarly, cluster 3 is the co-simulation cluster, showing that according to our taxonomy, co-simulation approaches are similar and have original features not found in the other categories. Cluster 2 contains coordination languages and ADLs. Even if the maximum distance between approaches is reduced to be considered part of the same cluster (clustering parameter), one cannot separate the approaches into two categories. Thus, we conclude that coordination languages and ADLs have more similarities in features compared to, for instance, co-simulation approaches or coordination frameworks.

In the end, Ptolemy remains unclustered, but we would classify Ptolemy as a coordination framework. However, it includes features from all categories and is not part of any cluster. It is fascinating that one approach can offer such diverse features.

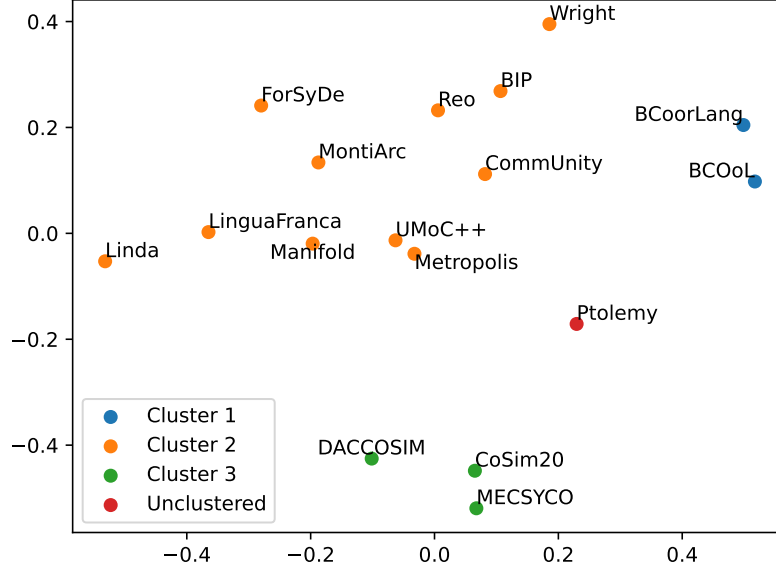


Fig. 7. Feature distance and clustering of approaches

7 Conclusion & future work

The main contribution of this paper is a taxonomy of coordination approaches represented as a feature model. This taxonomy facilitates the categorization and comparison of approaches across previously isolated categories. Consequently, we provide a comprehensive overview of coordination in the literature and the state of the art. Furthermore, we apply our taxonomy to 17 coordination approaches and publish the results as open data [51]. Analyzing the resulting data, we identify characteristic features in each category, provide overarching insights, and cluster the approaches regarding their similarity.

In future work, we aim to refine our taxonomy further and apply it to more approaches to obtain an even more comprehensive overview of coordination approaches.

References

1. Allen, R.: A Formal Approach to Software Architecture. Ph.D. thesis, Carnegie Mellon, School of Computer Science (Jan 1997)
2. Allen, R., Garlan, D.: A formal basis for architectural connection. *ACM Transactions on Software Engineering and Methodology* **6**(3), 213–249 (Jul 1997). <https://doi.org/10.1145/258077.258078>

3. André, C.: Syntax and semantics of the clock constraint specification language (CCSL). Technical Report RR-6925, INRIA (2009)
4. Arbab, F., Herman, I., Spilling, P.: An overview of manifold and its implementation. *Concurrency: Practice and Experience* **5**(1), 23–70 (Feb 1993). <https://doi.org/10.1002/cpe.4330050103>
5. Arbab, F.: What do you mean, coordination? *Bulletin of the Dutch Association for Theoretical Computer Science* pp. 11–22 (1998)
6. Arbab, F.: Reo: A channel-based coordination model for component composition. *Mathematical Structures in Computer Science* **14**(3), 329–366 (Jun 2004). <https://doi.org/10.1017/S0960129504004153>
7. Balarin, F., Watanabe, Y., Hsieh, H., Lavagno, L., Passerone, C., Sangiovanni-Vincentelli, A.: Metropolis: An integrated electronic system design environment. *Computer* **36**(4), 45–52 (Apr 2003). <https://doi.org/10.1109/MC.2003.1193228>
8. Basu, A., Bensalem, B., Bozga, M., Combaz, J., Jaber, M., Nguyen, T.H., Sifakis, J.: Rigorous Component-Based System Design Using the BIP Framework. *IEEE Software* **28**(3), 41–48 (May 2011). <https://doi.org/10.1109/MS.2011.27>
9. Bliudze, S., Sifakis, J.: The Algebra of Connectors—Structuring Interaction in BIP. *IEEE Transactions on Computers* **57**(10), 1315–1330 (Oct 2008). <https://doi.org/10.1109/TC.2008.26>
10. Camus, B., Galtier, V., Caujolle, M.: Hybrid co-simulation of FMUs using DEV&DESS in MECSYCO. In: 2016 Symposium on Theory of Modeling and Simulation (TMS-DEVS). pp. 1–8 (2016). <https://doi.org/10.23919/TMS.2016.7918814>
11. Camus, B., Paris, T., Vaubourg, J., Presse, Y., Bourjot, C., Ciarletta, L., Chevrier, V.: Co-simulation of cyber-physical systems using a DEVS wrapping strategy in the MECSYCO middleware. *SIMULATION* **94**(12), 1099–1127 (Dec 2018). <https://doi.org/10.1177/0037549717749014>
12. Carriero, N., Gelernter, D.: Linda in context. *Communications of The Acm* **32**(4), 444–458 (Apr 1989). <https://doi.org/10.1145/63334.63337>
13. Carriero, N.J., Gelernter, D., Mattson, T.G., Sherman, A.H.: The Linda alternative to message-passing systems. *Parallel Computing* **20**(4), 633–655 (Apr 1994). [https://doi.org/10.1016/0167-8191\(94\)90032-9](https://doi.org/10.1016/0167-8191(94)90032-9)
14. Clements, P.: A survey of architecture description languages. In: *Proceedings of the 8th International Workshop on Software Specification and Design*. pp. 16–25. IEEE Comput. Soc. Press, Schloss Velen, Germany (1996). <https://doi.org/10.1109/IWSSD.1996.501143>
15. Dad, C., Tavella, J.P., Vialle, S.: Synthesis and feedback on the distribution and parallelization of FMI-CS-based co-simulations with the DACCOSIM platform. *Parallel Computing* **106**, 102802 (Sep 2021). <https://doi.org/10.1016/j.parco.2021.102802>
16. Eker, J., Janneck, J., Lee, E., Jie Liu, Xiaojun Liu, Ludvig, J., Neuen-dorffer, S., Sachs, S., Yuhong Xiong: Taming heterogeneity - the Ptolemy approach. *Proceedings of the IEEE* **91**(1), 127–144 (Jan 2003). <https://doi.org/10.1109/JPROC.2002.805829>
17. Fiadeiro, J.L., Lopes, A.: Semantics of architectural connectors. In: Goos, G., Hartmanis, J., Van Leeuwen, J., Bidoit, M., Dauchet, M. (eds.) *TAPSOFT '97: Theory and Practice of Software Development*, vol. 1214, pp. 503–519. Springer Berlin Heidelberg, Berlin, Heidelberg (1997). <https://doi.org/10.1007/BFb0030622>
18. Galtier, V., Vialle, S., Dad, C., Tavella, J.P., Lam-Yee-Mui, J.P., Plessis, G.: FMI-Based distributed multi-simulation with DACCOSIM. In: *Proceedings of the Symposium on Theory of Modeling & Simulation: DEVS Integrative M&S Symposium*.

- pp. 39–46. DEVS '15, Society for Computer Simulation International, San Diego, CA, USA (2015)
19. Gomes, C., Thule, C., Broman, D., Larsen, P.G., Vangheluwe, H.: Co-Simulation: A Survey. *ACM Computing Surveys* **51**(3), 1–33 (May 2019). <https://doi.org/10.1145/3179993>
 20. Goos, G., Hartmanis, J., Van Leeuwen, J. (eds.): *Coordination Models and Languages*, vol. 2039, pp. 33–52. Springer Berlin Heidelberg, Berlin, Heidelberg (2001). https://doi.org/10.1007/3-540-44933-7_3
 21. Haber, A., Ringert, J.O., Rumpe, B.: MontiArc - Architectural Modeling of Interactive Distributed and Cyber-Physical Systems (2014). <https://doi.org/10.48550/ARXIV.1409.6578>
 22. Hafner, I., Popper, N.: An Overview of the State of the Art in Co-Simulation and Related Methods. *SNE Simulation Notes Europe* **31**(4), 185–200 (Dec 2021). <https://doi.org/10.11128/sne.31.on.10582>
 23. Hussain, S.: Investigating Architecture Description Languages (ADLs) A Systematic Literature Review (2013)
 24. Kang, K.C., Cohen, S.G., Hess, J.A., Novak, W.E., Peterson, A.S.: Feature-Oriented Domain Analysis (FODA) Feasibility Study:. Tech. rep., Defense Technical Information Center, Fort Belvoir, VA (Nov 1990). <https://doi.org/10.21236/ADA235785>
 25. Krauter, T.: Towards behavioral consistency in heterogeneous modeling scenarios. In: 2021 ACM/IEEE International Conference on Model Driven Engineering Languages and Systems Companion (MODELS-C). pp. 666–671. IEEE, Fukuoka, Japan (Oct 2021). <https://doi.org/10.1109/MODELS-C53483.2021.00107>
 26. Kräuter, T., König, H., Rutle, A., Lamo, Y., Stünkel, P.: Behavioral consistency in multi-modeling. *The Journal of Object Technology* **22**(2), 2:1 (2023). <https://doi.org/10.5381/jot.2023.22.2.a9>
 27. Le Coënt, A., Alexandre Dit Sandretto, J., Chapoutot, A.: Guaranteed cosimulation of cyber-physical systems (Mar 2020)
 28. Levandowsky, M., Winter, D.: Distance between Sets. *Nature* **234**(5323), 34–35 (Nov 1971). <https://doi.org/10.1038/234034a0>
 29. Liboni, G.: Complex Systems Co-Simulation with the CoSim20 Framework : For Efficient and Accurate Distributed Co-Simulations. Theses, Université Côte d’Azur (Apr 2021)
 30. Lohstroh, M.: Reactors: A Deterministic Model of Concurrent Computation for Reactive Systems. Ph.D. thesis, EECS Department, University of California, Berkeley (Dec 2020)
 31. Lohstroh, M., Menard, C., Bateni, S., Lee, E.A.: Toward a Lingua Franca for Deterministic Concurrent Systems. *ACM Transactions on Embedded Computing Systems* **20**(4), 1–27 (Jul 2021). <https://doi.org/10.1145/3448128>
 32. Magee, J., Dulay, N., Eisenbach, S., Kramer, J.: Specifying distributed software architectures. In: Goos, G., Hartmanis, J., Van Leeuwen, J., Schäfer, W., Botella, P. (eds.) *Software Engineering — ESEC '95*, vol. 989, pp. 137–153. Springer Berlin Heidelberg, Berlin, Heidelberg (1995). https://doi.org/10.1007/3-540-60406-5_12
 33. Malavolta, I., Lago, P., Muccini, H., Pelliccione, P., Tang, A.: What Industry Needs from Architectural Languages: A Survey. *IEEE Transactions on Software Engineering* **39**(6), 869–891 (Jun 2013). <https://doi.org/10.1109/TSE.2012.74>
 34. Mathaikutty, D.A., Patel, H.D., Shukla, S.K., Jantsch, A.: UMoC++: A C++-Based Multi-MoC Modeling Environment. In: Vachoux, A. (ed.) *Applications of Specification and Design Languages for SoCs*, pp. 115–130. Springer Netherlands, Dordrecht (2006). https://doi.org/10.1007/978-1-4020-4998-9_7

35. Medvidovic, N.: Moving Architectural Description from Under the Technology Lamppost. In: 32nd EUROMICRO Conference on Software Engineering and Advanced Applications (EUROMICRO'06). pp. 2–3. IEEE, Cavtat, Dubrovnik, Croatia (2006). <https://doi.org/10.1109/EUROMICRO.2006.47>
36. Medvidovic, N., Taylor, R.: A classification and comparison framework for software architecture description languages. *IEEE Transactions on Software Engineering* **26**(1), 70–93 (Jan/2000). <https://doi.org/10.1109/32.825767>
37. Medvidovic, N., Taylor, R.N.: A framework for classifying and comparing architecture description languages. In: Goos, G., Hartmanis, J., van Leeuwen, J., Jazayeri, M., Schauer, H. (eds.) *Software Engineering — ESEC/FSE'97*, vol. 1301, pp. 60–76. Springer Berlin Heidelberg, Berlin, Heidelberg (1997). https://doi.org/10.1007/3-540-63531-9_7
38. Nixon, L.J.B., Simperl, E., Krummenacher, R., Martin-Recuerda, F.: Tuplespace-based computing for the Semantic Web: A survey of the state-of-the-art. *The Knowledge Engineering Review* **23**(2), 181–212 (Jun 2008). <https://doi.org/10.1017/S0269888907001221>
39. Oliveira, C., Wermelinger, M.: The CommUnity Workbench. *Science of Computer Programming* **69**(1-3), 46–55 (Dec 2007). <https://doi.org/10.1016/j.scico.2006.09.005>
40. Omicini, A., Viroli, M.: Coordination models and languages: From parallel computing to self-organisation. *The Knowledge Engineering Review* **26**(1), 53–59 (Feb 2011). <https://doi.org/10.1017/S026988891000041X>
41. Ozkaya, M., Kloukinas, C.: Are We There Yet? Analyzing Architecture Description Languages for Formal Analysis, Usability, and Realizability. In: 2013 39th Euromicro Conference on Software Engineering and Advanced Applications. pp. 177–184. IEEE, Santander, Spain (Sep 2013). <https://doi.org/10.1109/SEAA.2013.34>
42. Pandey, R.K.: Architectural description languages (ADLs) vs UML: A review. *ACM SIGSOFT Software Engineering Notes* **35**(3), 1–5 (May 2010). <https://doi.org/10.1145/1764810.1764828>
43. Papadopoulos, G.A., Arbab, F.: Coordination Models and Languages. In: *Advances in Computers*, vol. 46, pp. 329–400. Elsevier (1998). [https://doi.org/10.1016/S0065-2458\(08\)60208-9](https://doi.org/10.1016/S0065-2458(08)60208-9)
44. Papadopoulos, G.A., Arbab, F.: Modelling activities in information systems using the coordination language MANIFOLD. In: *Proceedings of the 1998 ACM Symposium on Applied Computing - SAC '98*. pp. 185–193. ACM Press, Atlanta, Georgia, United States (1998). <https://doi.org/10.1145/330560.330667>
45. Ptolemaeus, C. (ed.): *System Design, Modeling, and Simulation: Using Ptolemy II*. UC Berkeley EECS Dept, Berkeley, Calif, 1. ed., version 1.02 edn. (2014)
46. Rensink, A.: The GROOVE Simulator: A Tool for State Space Generation. In: Kanade, T., Kittler, J., Kleinberg, J.M., Mattern, F., Mitchell, J.C., Naor, M., Nierstrasz, O., Pandu Rangan, C., Steffen, B., Sudan, M., Terzopoulos, D., Tygar, D., Vardi, M.Y., Weikum, G., Pfaltz, J.L., Nagl, M., Böhlen, B. (eds.) *Applications of Graph Transformations with Industrial Relevance*, vol. 3062, pp. 479–485. Springer Berlin Heidelberg, Berlin, Heidelberg (2004). https://doi.org/10.1007/978-3-540-25959-6_40
47. Rossi, D., Cabri, G., Denti, E.: Tuple-based Technologies for Coordination. In: Omicini, A., Zambonelli, F., Klusch, M., Tolksdorf, R. (eds.) *Coordination of Internet Agents*, pp. 83–109. Springer Berlin Heidelberg, Berlin, Heidelberg (2001). https://doi.org/10.1007/978-3-662-04401-8_4

48. Sander, I., Jantsch, A.: System Modeling and Transformational Design Refinement in ForSyDe. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems* **23**(1), 17–32 (Jan 2004). <https://doi.org/10.1109/TCAD.2003.819898>
49. Sander, I., Jantsch, A., Attarzadeh-Niaki, S.H.: ForSyDe: System Design Using a Functional Language and Models of Computation. In: Ha, S., Teich, J. (eds.) *Handbook of Hardware/Software Codesign*, pp. 1–42. Springer Netherlands, Dordrecht (2016). https://doi.org/10.1007/978-94-017-7358-4_5-1
50. Schweiger, G., Gomes, C., Engel, G., Hafner, I., Schoeggl, J., Posch, A., Nouidui, T.: An empirical survey on co-simulation: Promising standards, challenges and research needs. *Simulation Modelling Practice and Theory* **95**, 148–163 (Sep 2019). <https://doi.org/10.1016/j.simpat.2019.05.001>
51. Tim Kräuter: Artifacts Coordination. <https://github.com/timKraeuter/Coordination-2024> (Feb 2024)
52. Vara Larsen, M.: BCOol : The Behavioral Coordination Operator Language. Ph.D. thesis, Université Nice Sophia Antipolis (Apr 2016)
53. Vara Larsen, M.E., Deantoni, J., Combemale, B., Mallet, F.: A Behavioral Coordination Operator Language (BCOoL). In: 2015 ACM/IEEE 18th International Conference on Model Driven Engineering Languages and Systems (MODELS). pp. 186–195. IEEE, Ottawa, ON, Canada (Sep 2015). <https://doi.org/10.1109/MODELS.2015.7338249>
54. Woods, E., Hilliard, R.: Architecture Description Languages in Practice Session Report. In: 5th Working IEEE/IFIP Conference on Software Architecture (WICSA'05). pp. 243–246. IEEE, Pittsburgh, PA, USA (2005). <https://doi.org/10.1109/WICSA.2005.15>